




# Feature Selection on Imbalanced Data

Presented by Randall Wald



# Introduction


- Data mining background
- What is feature selection?
- Why feature selection matters
- Traditional approaches
- Working with imbalanced data





# Data Mining

- Many instances
  - ◆ Each has features
  - ◆ Each has a class
- The goal: Create a classifier based on the features which predicts the class






# Numeric vs. Nominal

- Features and classes can each be numeric or nominal
- Numeric classes have certain classifiers:
  - ◆ Regressions
  - ◆ Module Order Modeling
- Nominal classes: good/bad, red/green, categories
- Features also are numeric or nominal



# Classifier: Naïve Bayes


- For each feature, what fraction of instances of each value are in each class?
- When evaluating instances, add together the probabilities of being in each class
- Assign to class based on probability





# Classifier: Decision Tree

- Find feature which best predicts which class each instance is in
- Make a flow chart: instances start at the top, and go to a different child node based on what value they have for that feature
- Iterate to create further children





# Classifier: Nearest Neighbor


- Decide upon a decision metric to say how far apart two instances are
  - ◆ Geometric distance
  - ◆ Manhattan distance
- For new instances, find to nearest neighbors in each class
- Assign to class with nearest neighbor





# Classifier: SVMs

- Support Vector Machines
- Plot data in n-dimensional space
- Find hyperplane which best separates the data







# Feature Selection


- Many real-world data sets have more features than instances:
  - ◆ Gene chip data
  - ◆ Risk assessment
  - ◆ Text mining
  - ◆ Performance evaluation





# Gene chips

- Thousands of wells
- Each a separate gene to test
- One sample is run against all genes
- Sample = instance
- Genes = features
- Chips are expensive, so few samples





# Risk assessment

- Dangerous vs. not dangerous
  - ◆ Terrorism
  - ◆ Security
  - ◆ Financial
- Many factor to check
  - ◆ Buying habits
  - ◆ Movement patterns
  - ◆ Communication
- Few samples to work with



# Text mining


- Each distinct word is a feature
  - ◆ Only two states: present and absent
- Relatively few documents, at least compared to the number of total words
- Goal: Use words to determine which documents are in various categories
- Applications to web search, etc.





# Performance evaluation

- Predicting the quality of future systems based on existing systems
- Many properties to describe a system
  - ◆ Hardware specs
  - ◆ Software properties
- Few systems to compare when planing new systems






# Feature Selection

- Many algorithms become computationally intensive
  - ◆ Naïve Bayes is  $O(n)$
  - ◆ Decision Tree is  $O(n^2)$  for complete tree
- Extraneous, noisy features drown out noise
  - ◆ That many features, some random ones will have a pattern
  - ◆ Hard to decide weightings



# Feature Selection


- Solution: Feature selection
- Pick some subset of the features
- Only run the classifier on those features
- Results are often similar to using all features





# Choosing features

- Expert knowledge
- Filter
  - ◆ Feature ranking
  - ◆ Subset evaluation
- Wrapper
- Hybrid







# Expert knowledge

- Some attributes are known to matter more than others
- Experts can eliminate obviously-useless features
- Also aids identification of linked features in advance





# Filter: Feature Ranking

- Perform evaluation on each feature
  - ◆ Information gain
  - ◆ Odds ratio
  - ◆ Chi-squared
  - ◆ OneR
- Rank features by performance
- Select the top few features
  - ◆ Number to use found via statistical measure or pilot study



# Information gain


- Amount of entropy the data have with and without the feature in question
- Shannon entropy: How many bits it would take to encode the data in question
- Intuitively, how much we know
- If removing a feature removes information, it's a good feature





# Information gain


- Calculation of entropy from a given state:  $e(a, b, c, \dots) = -\log(a/n) - \log(b/n) - \log(c/n) \dots$
- $a$  is number of instances in class  $a$ ,  $b$  is number of instances in class  $b$ , etc.
- $n = a + b + c \dots =$  total number of instances in current state





# Information gain


- Calculation of entropy for a given feature:
  - ◆ For each value of that feature, calculate entropy (that is, use the earlier equation for instances with that value for that feature)
  - ◆ Add up entropies for all values, weighted by number of instances with that value





# Information gain


- Information gain for a given feature
  - ◆ Calculate entropy of initial state (that is, the entire set)
  - ◆ Information gain = total entropy – entropy of that feature
- Rank features by information gain, higher is better





# Odds ratio

- Only works with binary features (those which have just two values, present and absent) and binary classes (true and false)
- Ratio of the odds of a feature being present in the true instance versus those of it being in the false instances





# Odds ratio

- Odds of being present in true instances  
= probability of being in true instances /  
probability of not being in true instance  
=  $\text{tpr} / (1 - \text{tpr})$
- Odds of being present in false instances  
=  $\text{fpr} / (1 - \text{fpr})$
- Odds ratio =  $\frac{\text{tpr} (1 - \text{fpr})}{(1 - \text{tpr}) \text{fpr}}$





# Chi-squared

- Assume that feature distribution is unrelated to class distribution
- Calculate the expected distribution of feature values
- Determine how far away from this expectation the actual feature distribution is
- Higher chi-squared is better



# OneR metric

- Extremely simple classifier
  - ◆ For each value, predict the class which predominates amongst instances with that value
  - ◆ For numeric attributes, break into intervals
- Classify all instances
- Evaluate accuracy of classification
- Better classification = better feature

# When to stop selection?

- Based on the statistical measure
  - ◆ Information Gain, Odds Ratio, etc. get too low to be useful
- Based on ad-hoc number
  - ◆ We only want the top 20 features
- Pilot study
  - ◆ On a smaller sample, try 5, 10, ..., 75 features, and measure the performance of the resulting classifiers



# Filter: Subset evaluation

- Select subsets to examine
  - ◆ Exhaustive search
  - ◆ Greedy stepwise
  - ◆ Best first
- Perform evaluation on each subset
  - ◆ Correlation-based feature selection
  - ◆ Consistency
  - ◆ Markov blanket
- Stop search when a threshold is met and use just those features



# Search techniques

- Exhaustive search
  - ◆ Just try all possible subsets
  - ◆ Impractical in most situations
- Greedy stepwise
  - ◆ Start with either empty or full set
  - ◆ Add or remove best or worst feature
- Best first
  - ◆ Start with empty set
  - ◆ Add best feature
  - ◆ Compare with previous sets; if new feature doesn't help, backtrack



# Correlation (CFS)

- Two competing goals
  - ◆ Features which predict (correlate) highly with the class
  - ◆ Features which do not correlate highly with one another
- Metric to evaluate performance:  
Pearson Correlation Coefficient



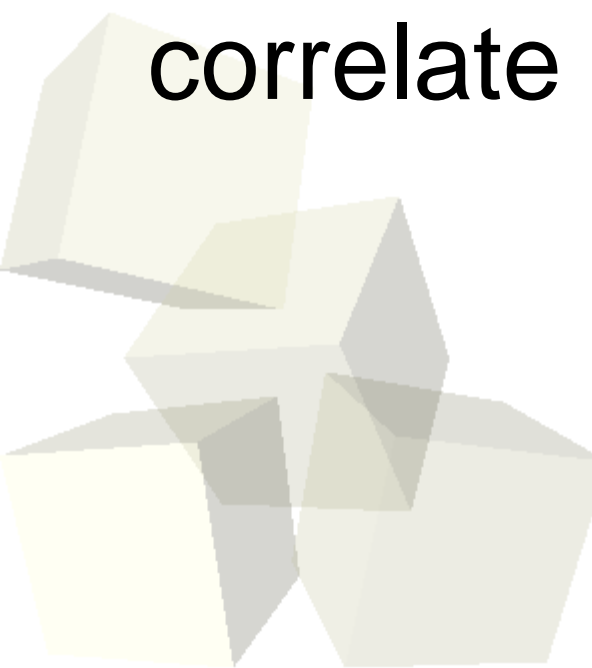
# Pearson correlation

- For two variables, the variance in Y which is accounted for by the variance of X
- For k variables. more complex:

$$M_S = \frac{k\overline{r_{cf}}}{\sqrt{k + k(k-1)\overline{r_{ff}}}}$$

- $M_S$  is merit,  $r_{cf}$  is correlation between features and class, and  $r_{ff}$  is pairwise correlation between features

- Actual methods of calculating the  $r_{cf}$  and  $r_{ff}$  values vary
- Numerator is how well features predict class
- Denominator is how much features correlate with one another


$$M_S = \frac{k\overline{r_{cf}}}{\sqrt{k + k(k-1)\overline{r_{ff}}}}$$





# Consistency

- Only works with random or exhaustive search
- Generate new set
- If it's smaller than the current best
  - ◆ See if subset is sufficiently consistent
  - ◆ If it is, it's the new best subset
  - ◆ Print this subset
- If it's identical in size to the best
  - ◆ See if subset is sufficiently consistent
  - ◆ Print this subset



# Consistency measure

- Want subsets that “make sense” as descriptions of the class
- Subsets are inconsistent if two instances with identical features have different classes
- For each set of *instances* with matching feature values, find how many are inconsistent
- $\frac{\text{\# of inconsistent instances}}{\text{total \# of instances}} = \text{inconsistency}$



# Markov blanket

- Based on conditional probabilities: probability of being in a given class based on what we know about the features
- Want to remove features which do not add any additional information to the current feature set






# Markov blanket

- A feature  $F_i$  is “conditionally independent” of some other subset  $C$  of features if the conditional probabilities for  $C + F_i$  are identical to those for  $C$  alone
- In this case,  $F_i$  is an unnecessary feature, since it adds no new information



# Markov blanket

- In practice, finding conditional probabilities for large sets of features is hard
- Instead, test with many small subsets of features





# Markov blanket

- Start with full set of features
- For each feature  $F_i$ , find feature subset of size  $k$  which best correlates with  $F_i$  (not including  $F_i$  itself)
- Find how conditionally independent  $F_i$  is from that subset
- After trying all  $F_i$ , remove the worst one and iterate on those remaining




# Wrapper

- Select subsets (same as subset evaluation)
- Use a classifier with just those features; test performance
- Performance of classifier on subset is quality of subset
- Stop search when a threshold is met



# Wrapper

- “Wrapped” classifier is usually identical to the target classifier, but not always
- Different performance metrics can rank the subsets







# Comparisons

- Feature ranking is much cheaper than subset evaluation or wrapper
- Subset evaluation and wrapper avoid the chance that two highly-ranked features don't add new information together
- Wrapper is tightly bound to the target classifier, which is good and bad



# Hybrid approaches


- Some feature selection methods excel at removing completely useless features
- Others are better at comparing among important features to determine most important
- Solution: apply more than one filter





# Hybrid approaches


- Especially useful for using a feature ranker to filter out the really bad features, then subset selection or a wrapper on the remainder
- Example: Odds ratio followed by Markov blanket





# Feature Selection

- Not all approaches work with all classifiers
- Many parameters that need to be varied to maximize performance
- Overfitting a perennial risk





# Imbalanced data

- Frequently interested in binary classes
  - ◆ Sick/healthy
  - ◆ Risky/safe
  - ◆ Fault-prone/not fault-prone
- Many more examples of one class than the other
- Care more about the minority class





# Imbalanced data

- Many classifiers seek to balance Type I and Type II errors
  - ◆ Type I: false positives
  - ◆ Type II: false negatives
- For imbalanced data, this causes all positive instances to be classified as negative instances
  - ◆ Even if all positive instances classified as negative, still small compared to any false positives



# Imbalanced data

- Some classifiers can be modified to handle this
  - ◆ Naïve Bayes: require lower probability to find positive class
  - ◆ Nearest Neighbor: give positive class multiplier on distance
- These approaches require fine-tuning





# Imbalanced data

- The same data sets which have many features often have imbalanced data
- Feature selection techniques also have problems on such data
  - ◆ Information Gain and Odds Ratio favor negative features
  - ◆ Subset evaluation and wrapper usually favor subsets which balance Type I and Type II error






# Imbalanced data

## ■ Solutions:


- ◆ Use different feature ranking metrics
  - Bi-normal separation
  - FAST
- ◆ Sampling
  - Effectiveness varies based on feature selection and classifier





# Bi-normal separation

- BNS: feature ranking method designed for imbalanced data
- Only works for binary features
- Thus, useful for text mining and evaluating classifiers





# Bi-normal separation

- Find # of false positives, # of true positives, total # of positives, total # of negatives
- $\text{\#tp} / \text{total positives} = \text{true positive rate}$
- $\text{\#fp} / \text{total negatives} = \text{false positive rate}$
- Find inverse normal CDF for fpr and tpr
- Difference is BNS



# Bi-normal separation



**Figure 1. Two views of Bi-Normal Separation using the Normal probability distribution: (left) Separation of thresholds. (right) Separation of curves (ROC analysis).**

- Assumes that fpr and tpr are fixed values
- Probability of feature being present is normally distributed



# Bi-normal separation




**Figure 1. Two views of Bi-Normal Separation using the Normal probability distribution: (left) Separation of thresholds. (right) Separation of curves (ROC analysis).**

- Difference between tpr and fpr is real classification power of feature
- Selects both positive and negative features




# FAST

- Feature Assessment by Sliding Thresholds: approximate area under ROC curve for threshold-based classifiers run on each feature






# ROC curves



- tpr vs fpr
- (0,0) corresponds to always guessing negative
- (1,1) is always guessing positive
- $x=y$  line is random guess



# ROC curves




- ROC works on ranges of classifiers
  - ◆ Naïve Bayes with sliding thresholds
  - ◆ OneR with different thresholds





# ROC curves



- Area under ROC curve corresponds to “goodness” of models
- B is better than A in this example




# FAST

- Create a collection of models for each feature:
  - ◆ 10 threshold models
  - ◆ Each threshold is the mean of one quantile
- Find ROC for each feature
  - ◆ If you get an ROC of  $< 0.5$ , take the value  $(1 - \text{ROC})$
- ROC value is quality of feature



# FAST

- Why use ROC and sliding thresholds?
  - ◆ Hard to pick one good static threshold
  - ◆ ROC is the final performance metric, so might as well use for feature selection
- FAST only uses 10 thresholds, but should approximate the use-all-possible-thresholds case






# Sampling

- Modify the data set to not be imbalanced anymore
  - ◆ Oversampling: duplicate instances of positive class
  - ◆ Undersampling: remove instances of the negative class
- Randomized sampling, or focused on strengthening the boundary
  - ◆ Oversample near boundary
  - ◆ Undersample far from boundary



# Sampling

- Oversampling can be applied both to feature selection and to building a classifier with those features
- Efficacy of oversampling depends on specific properties of feature selection or classifier





# Sampling

## ■ Feature selection

- ◆ OR, BNS: Immune to oversampling
- ◆  $\text{Chi}^2$ , IG: Less biased with sampling

## ■ Classifiers

- ◆ Decision Tree: Immune to sampling of features, susceptible to sampling of itself
- ◆ Naïve Bayes: Sampling of itself and features both help
- ◆ SVMs: Only feature sampling helps



# Future directions

- Wrapper with different performance metric
  - ◆ Instead of finding subsets with the best accuracy (I/II ratio), maximize BNS or ROC of the subset
- Novel imbalance-aware subset evaluator
- Combining existing approaches in new ways



# Conclusion

- Both feature selection and imbalanced data are important problems in data mining
- Solving both of them together is essential to creating classifiers for a variety of applications







# References

- X. Chen and M. Wasikowski, “FAST: A ROC-based Feature Selection Metric for Small Samples and Imbalanced Data Classification Problems,” *Proc. 14th ACM SIGKDD Int'l Conf. Knowledge Discovery & Data Mining* (KDD'08), 2008, pp. 124-132.
- G. Forman, “An Extensive Empirical Study of Feature Selection Metrics for Text Classification,” *J. Machine Learning Research*, vol. 3, May 2003, pp. 1289-1305.
- M. Hall, “Correlation-based Feature Selection for Machine Learning,” doctoral dissertation, Dept. Computer Science, Univ. of Waikato, 1999.
- R. C. Holte, “Very Simple Classification Rules Perform Well on Most Commonly Used Datasets,” *Machine Learning*, vol. 11, April 1993, pp. 63-91.
- D. Koller and M. Sahami, “Toward Optimal Feature Selection,” *Proc. 13<sup>th</sup> Int'l Conf. Machine Learning* (ICML'96), 1996.
- H. Liu and R. Setiono, “A Probabilistic Approach to Feature Selection – A Filter Solution,” *Proc. 13<sup>th</sup> Int'l Conf. Machine Learning* (ICML'96), 1996.
- E. Xing, M. Jordan, and R. Karp, “Feature Selection for High-Dimensional Geometric Microarray Data,” *Proc. 18<sup>th</sup> Int'l Conf. Machine Learning* (ICML'01), 2001.